

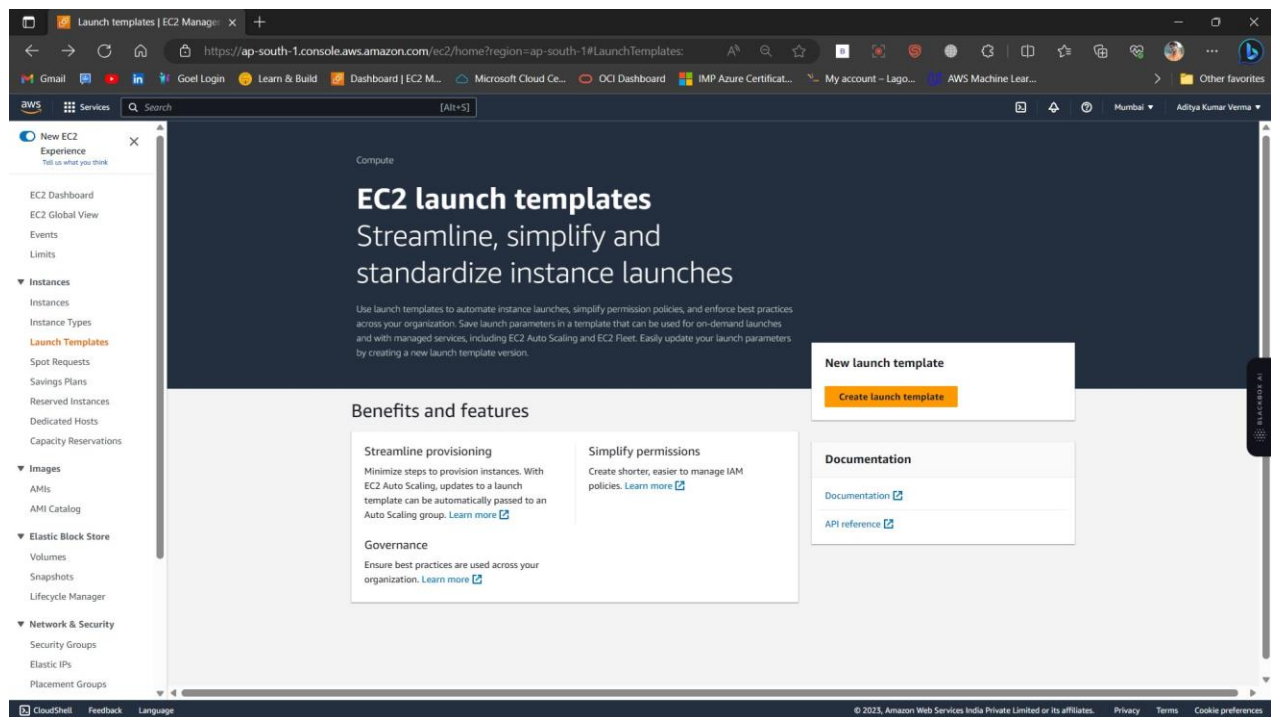
TASK 3(b)

Auto-Scaling as accordance with CPU-utilization:

Min: 2 Desired: 2 Maximum: 5

Step 1: Open you AWS Management Console.

Step 2: Go to Launch Template and Launch one.



Step 3: Give a name as “AdiTemplate”, and type “39” in version.

Create launch template

Creating a launch template allows you to create a saved instance configuration that can be reused, shared and launched at a later time. Templates can have multiple versions.

Launch template name and description

Launch template name - required
AdiTemplate
Must be unique to this account. Max 128 chars. No spaces or special characters like '%', '"', '@'.

Template version description
39
Max 255 chars

Auto Scaling guidance [Info](#)
Select this if you intend to use this template with EC2 Auto Scaling
☐ Provide guidance to help me set up a template that I can use with EC2 Auto Scaling

► Template tags
► Source template

Launch template contents
Specify the details of your launch template below. Leaving a field blank will result in the field not being included in the launch template.

▼ **Application and OS Images (Amazon Machine Image)** [Info](#)
An AMI is a template that contains the software configuration (operating system, application server, and applications) required to launch your instance. Search or Browse for AMIs if you don't see what you are looking for below

Search our full catalog including 1000s of application and OS images

Summary

Software Image (AMI)
-

Virtual server type (instance type)
-

Firewall (security group)
-

Storage (volumes)
-

Free tier: In your first year includes 750 hours of t2.micro (or t3.micro in the Regions in which t2.micro is unavailable) instance usage on free tier AMIs per month, 30 GiB of EBS storage, 2 million I/Os, 1 GiB of snapshots, and 100 GiB of bandwidth to the internet.

Cancel Create launch template

Step 4: Select “AWS Linux” for OS.

Launch template contents

Specify the details of your launch template below. Leaving a field blank will result in the field not being included in the launch template.

▼ **Application and OS Images (Amazon Machine Image)** [Info](#)
An AMI is a template that contains the software configuration (operating system, application server, and applications) required to launch your instance. Search or Browse for AMIs if you don't see what you are looking for below

Search our full catalog including 1000s of application and OS images

Quick Start

Don't include in launch template

Amazon Linux
macOS
Ubuntu
Windows
Red Hat

Browse more AMIs
Including AMIs from AWS, Marketplace and the Community

Amazon Machine Image (AMI)

Amazon Linux 2023 AMI
ami-0b9ec7f1fe947bbdd (64-bit (x86), uefi-preferred) / ami-0b9ec7f1fe947bbdd (64-bit (ARM), uefi)
Virtualization: hvm ENA enabled: true Root device type: ebs Free tier eligible

Description
Amazon Linux 2023 AMI 2023.1.20230629.0 x86_64 HVM kernel-5.1

Architecture
64-bit (x86)

Boot mode
uefi-preferred

AMI ID
ami-0b9ec7f1fe947bbdd Verified provider

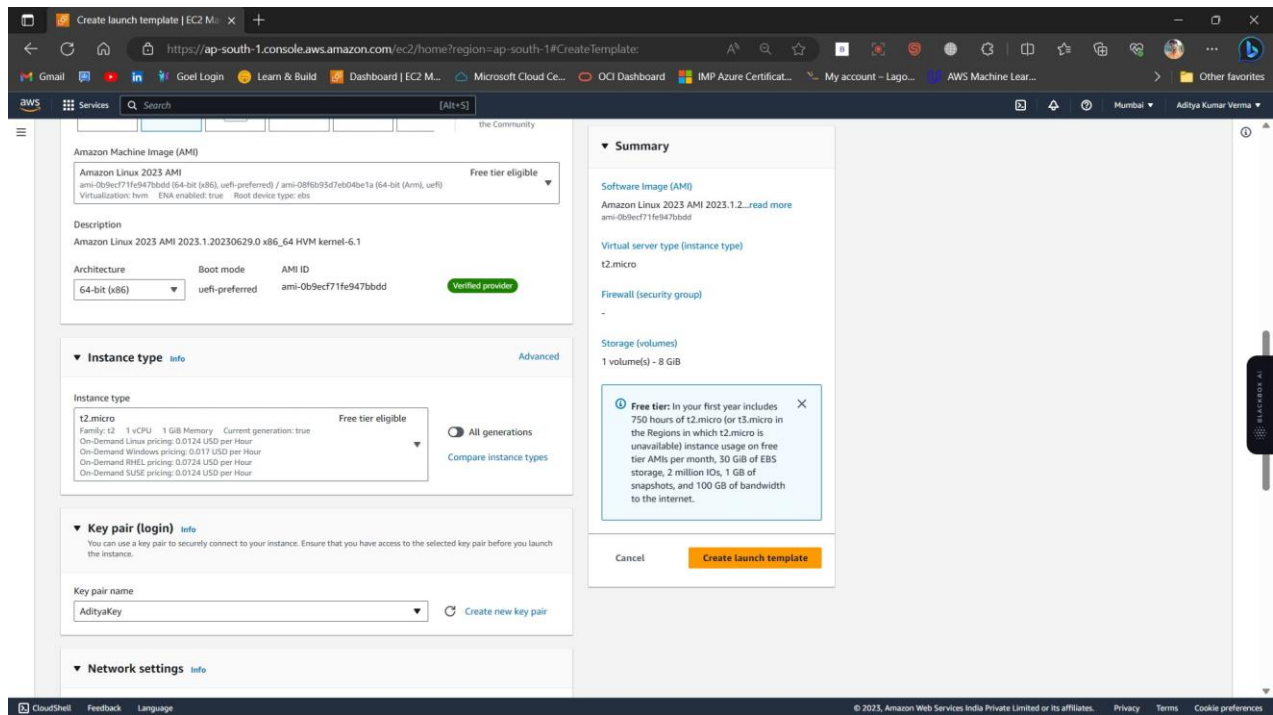
Instance type [Info](#) [Advanced](#)

Instance type
Don't include in launch template

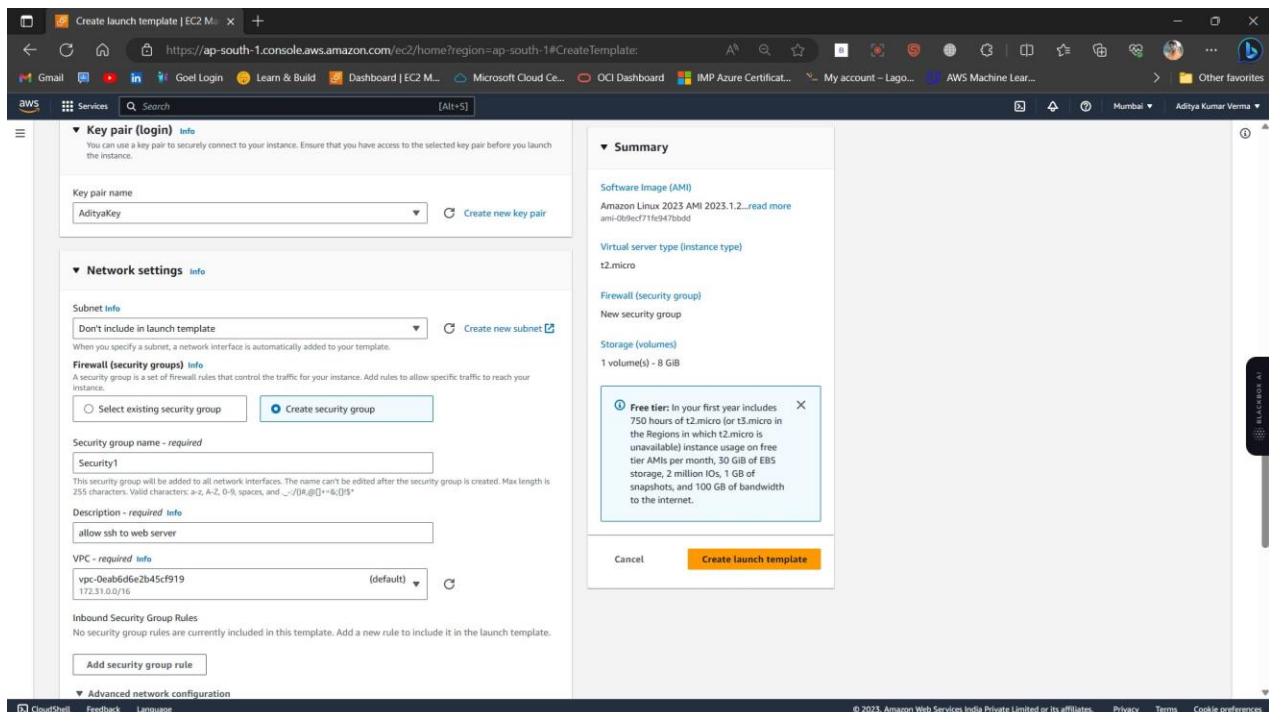
Free tier: In your first year includes 750 hours of t2.micro (or t3.micro in the Regions in which t2.micro is unavailable) instance usage on free tier AMIs per month, 30 GiB of EBS storage, 2 million I/Os, 1 GiB of snapshots, and 100 GiB of bandwidth to the internet.

Cancel Create launch template

Step 5: Select “t2 micro” at instance and select a keypair “AdityaKey”.



Step 6: Create a security group and give it a name as “Security1”
& add description.



Step 7: Add two inbound rules 1. HTTP and 2. SSH.

allow ssh to web server

VPC - required [info](#)
vpc-0eab6d6e2b45cf919 (default)

Inbound Security Group Rules

Security group rule 1 (TCP, 80) [Remove](#)

Type [info](#) HTTP Protocol [info](#) TCP Port range [info](#) 80

Source type [info](#) Custom Source [info](#) Add CIDR, prefix list or security Description - optional [info](#) e.g. SSH for admin desktop

Security group rule 2 (TCP, 22) [Remove](#)

Type [info](#) ssh Protocol [info](#) TCP Port range [info](#) 22

Source type [info](#) Custom Source [info](#) Add CIDR, prefix list or security Description - optional [info](#) e.g. SSH for admin desktop

Add security group rule

Advanced network configuration

No network interfaces are currently included in this template. Add a network interface to include it in the launch template.

Add network interface

Storage (volumes) [info](#)

1 volume(s) - 8 GiB

Free tier: In your first year includes 750 hours of t2.micro (or t3.micro in the Regions in which t2.micro is unavailable) instance usage on free tier AMIs per month, 30 GiB of EBS storage, 2 million I/Os, 1 GiB of snapshots, and 100 GiB of bandwidth to the internet.

Cancel Create launch template

Step 8: In additional settings add this script and click on “Create launchtemplate”.

Metadata [info](#)

Don't include in launch template

Metadata transport [info](#)
Don't include in launch template

Metadata version [info](#)
Don't include in launch template

Metadata response hop limit [info](#)
Don't include in launch template

Allow tags in metadata [info](#)
Don't include in launch template

User data - optional [info](#)
Upload a file with your user data or enter it in the field.

Choose file

```
sudo -i
yum update
yum install nginx
systemctl start nginx
```

☐ User data has already been base64 encoded

Summary

Software Image (AMI)
Amazon Linux 2023 AMI 2023.1.2...[read more](#)
ami-0b9ecf71fe947bbdd

Virtual server type (instance type)
t2.micro

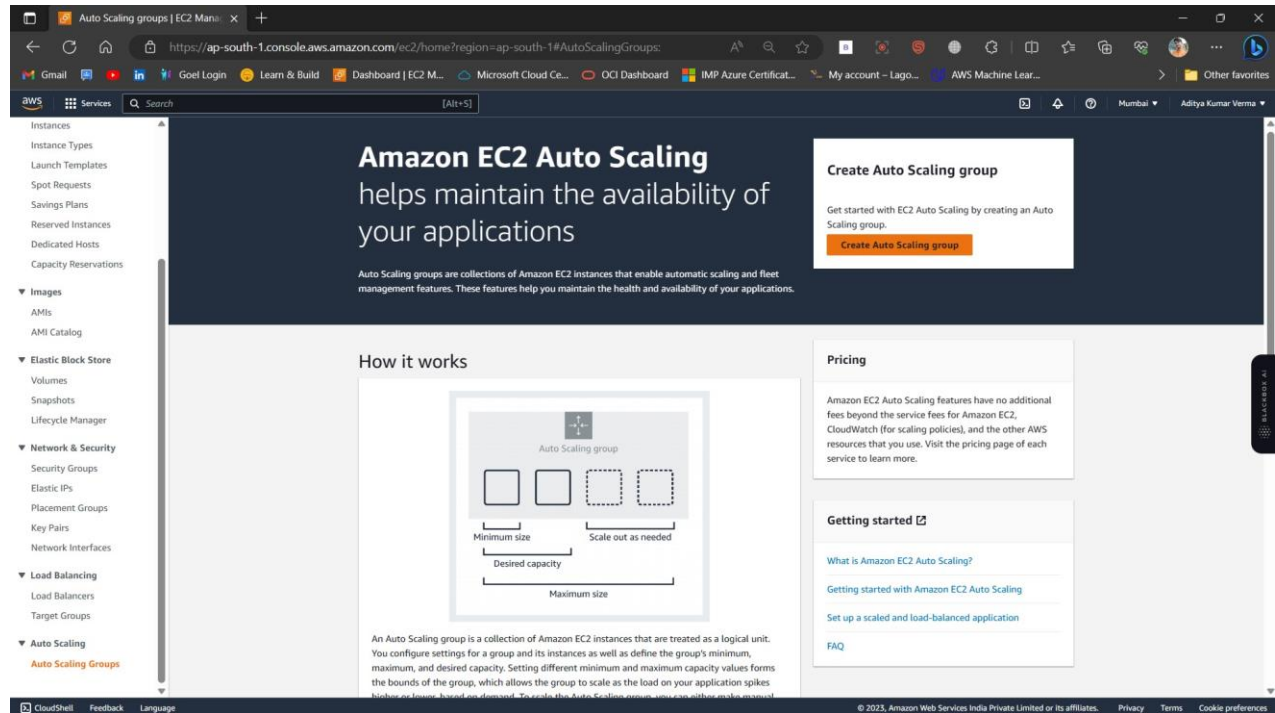
Firewall (security group)
New security group

Storage (volumes)
1 volume(s) - 8 GiB

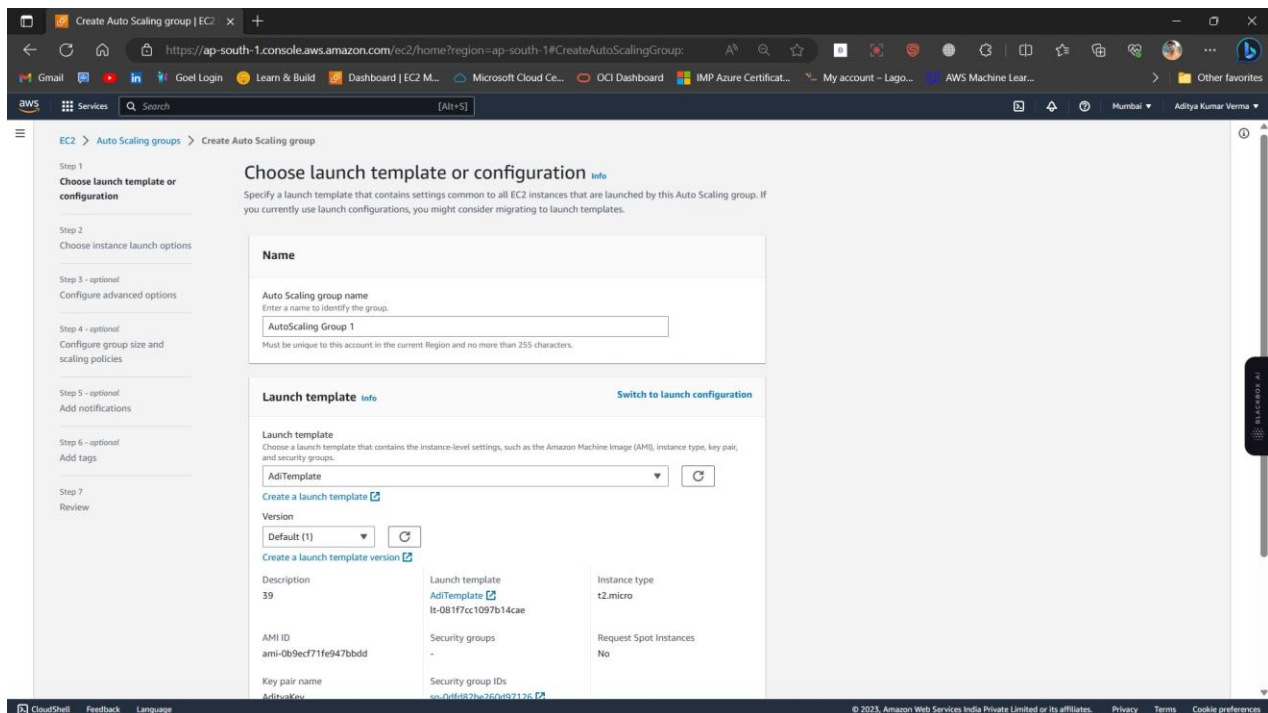
Free tier: In your first year includes 750 hours of t2.micro (or t3.micro in the Regions in which t2.micro is unavailable) instance usage on free tier AMIs per month, 30 GiB of EBS storage, 2 million I/Os, 1 GiB of snapshots, and 100 GiB of bandwidth to the internet.

Cancel Create launch template

Step 9: Now go to Auto-Scaling and click on “create Auto Scaling group”.



Step 10: Give a name as “AutoScaling Group 1” and select the template you just created.



Step 11: Select “ap-south 1a” and “ap-south 1b” zones and click next.

The screenshot shows the 'Create Auto Scaling group' wizard in the AWS Management Console, specifically Step 2: 'Choose instance launch options'. The left sidebar shows the progress: Step 1 (Choose launch template or configuration), Step 2 (Choose instance launch options), Step 3 (optional: Configure advanced options), Step 4 (optional: Configure group size and scaling policies), Step 5 (optional: Add notifications), Step 6 (optional: Add tags), and Step 7 (Review). The main content area is titled 'Choose instance launch options' and includes a 'Network' section. The 'Network' section has a 'VPC' dropdown set to 'vpc-0eab6d6e2b45cf919' and an 'Availability Zones and subnets' section with three selected subnets: 'ap-south-1a | subnet-09e262654103128b5', 'ap-south-1b | subnet-097d712c33580afe7', and 'ap-south-1c | subnet-0a7705bec02c3f4c'. Below this is an 'Instance type requirements' section with an 'Override launch template' button. The bottom of the page shows the AWS footer with copyright information and links for CloudShell, Feedback, Language, Privacy, Terms, and Cookie preferences.

Step 12: Reduce health period to “50” seconds .

The screenshot shows the 'Create Auto Scaling group' wizard in the AWS Management Console, specifically Step 6: 'Configure group size and scaling policies'. The left sidebar shows the progress: Step 5 (optional: Add notifications), Step 6 (optional: Add tags), and Step 7 (Review). The main content area is titled 'Configure group size and scaling policies' and includes a 'Health checks' section. The 'Health checks' section has a 'Health check grace period' dropdown set to '50' seconds. Below this is an 'Additional settings' section with a 'Monitoring' checkbox labeled 'Enable group metrics collection within CloudWatch' and a 'Default instance warmup' checkbox labeled 'Enable default instance warmup'. The bottom of the page shows the AWS footer with copyright information and links for CloudShell, Feedback, Language, Privacy, Terms, and Cookie preferences.

Step 13: Select Group size as ->minimum: 2; desired:2;
maximum:5

The screenshot shows the 'Configure group size and scaling policies' step of the AWS Auto Scaling group creation process. The interface includes a sidebar with steps 1 through 7, and a main content area with three sections: 'Group size - optional', 'Scaling policies - optional', and 'Instance scale-in protection - optional'.

Group size - optional

Specify the size of the Auto Scaling group by changing the desired capacity. You can also specify minimum and maximum capacity limits. Your desired capacity must be within the limit range.

Desired capacity: 2

Minimum capacity: 2

Maximum capacity: 5

Scaling policies - optional

Choose whether to use a scaling policy to dynamically resize your Auto Scaling group to meet changes in demand.

☐ Target tracking scaling policy
Choose a desired outcome and leave it to the scaling policy to add and remove capacity as needed to achieve that outcome.

☒ None

Instance scale-in protection - optional

Instance scale-in protection

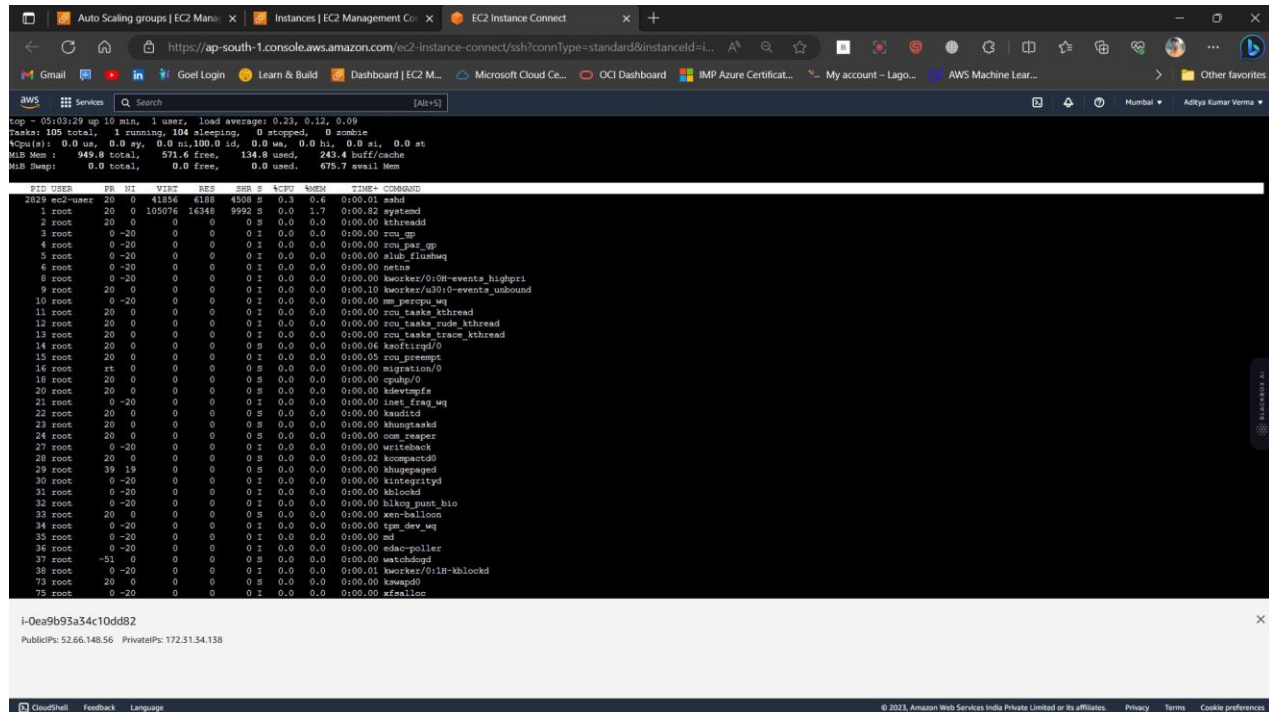
Step 14: Select “Target Tracking scaling policy”
& set Target Value at 30.

Go back to instances you will find two unnamed instances running.

The screenshot shows the 'Instances' page in the AWS Management Console. It displays a table of two running EC2 instances. The table has columns for Name, Instance ID, Instance state, Instance type, Status check, Alarm status, Availability Zone, Public IPv4 DNS, Public IPv4, Elastic IP, and IPv6 IPs.

Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 DNS	Public IPv4	Elastic IP	IPv6 IPs
-	i-0ea9b93a34c10d682	Running	t2.micro	2/2 checks passed	No alarms	ap-south-1a	ec2-52-66-148-56.ap-s...	52.66.148.56	-	-
-	i-0a1462c6d77543c87	Running	t2.micro	2/2 checks passed	No alarms	ap-south-1b	ec2-3-111-30-155.ap-s...	3.111.30.155	-	-

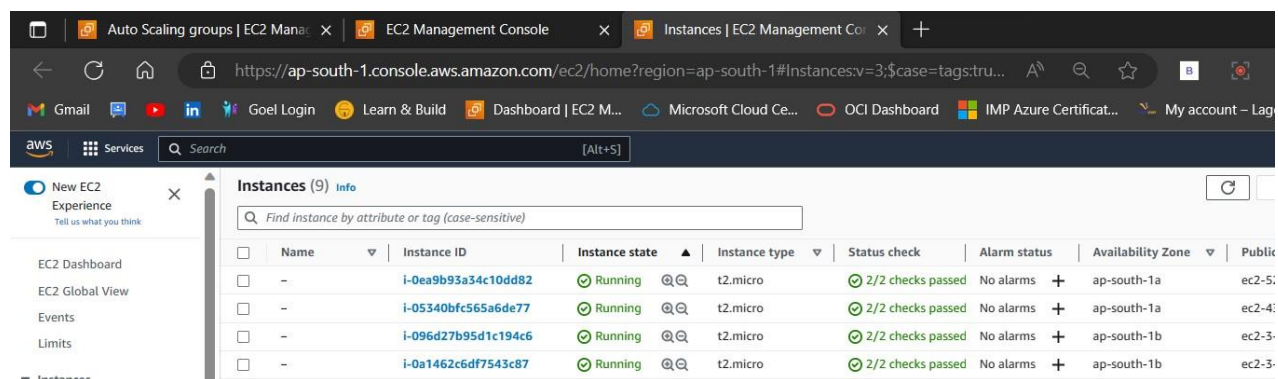
Step 15: Launch one of the instances and check the cpu capacity running by typing “top” command.



```
top - 01:03:29 up 10 min, 1 user, load averages: 0.23, 0.12, 0.09
tasks: 105 total, 1 running, 104 sleeping, 0 stopped, 0 zombie
%cpu(s): 0.0 us, 0.0 sy, 0.0 ni, 100.0 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
Mem Mem: 949.6 total, 571.6 free, 134.8 used, 243.4 buff/cache
Mem Swap: 0.0 total, 0.0 free, 0.0 used, 675.7 avail Mem

  PID USER      PR  NI  VIRT  RES  SHR S %CPU  %MEM    TIME+  COMMAND
 2029 root      20   0   112K   61K  4508 S  0.3   0.4   0:00.01  sshd
    1 root      20   0   105076 16348 9992 S  0.0   1.7   0:00.02  systemd
    2 root      20   0   0        0      0 S  0.0   0.0   0:00.00  kthreadd
    3 root      0 -20   0        0      0 S  0.1   0.0   0:00.00  rcu_gp
    4 root      0 -20   0        0      0 S  0.1   0.0   0:00.00  rcu_par_gp
    5 root      0 -20   0        0      0 S  0.1   0.0   0:00.00  slab_flushq
    6 root      0 -20   0        0      0 S  0.1   0.0   0:00.00  netns
    8 root      0 -20   0        0      0 S  0.1   0.0   0:00.00  kworker/0:0H-events_highpri
    9 root      20   0   0        0      0 S  0.1   0.0   0:00.10  kworker/u30:0-events_unbound
   10 root      0 -20   0        0      0 S  0.1   0.0   0:00.00  mm_percpu_wq
   11 root      20   0   0        0      0 S  0.1   0.0   0:00.00  rcu_tasks_kthreadd
   12 root      20   0   0        0      0 S  0.1   0.0   0:00.00  rcu_tasks_rude_kthreadd
   13 root      20   0   0        0      0 S  0.1   0.0   0:00.00  rcu_tasks_trace_kthreadd
   14 root      20   0   0        0      0 S  0.0   0.0   0:00.04  ksoftirqd/0
   15 root      20   0   0        0      0 S  0.0   0.0   0:00.05  rcu_preempt
   16 root      rt   0   0        0      0 S  0.0   0.0   0:00.00  migration/0
   18 root      20   0   0        0      0 S  0.0   0.0   0:00.00  cpuhp/0
   20 root      20   0   0        0      0 S  0.0   0.0   0:00.00  kdevtmpfs
   21 root      0 -20   0        0      0 S  0.1   0.0   0:00.00  inet_frag_wq
   22 root      20   0   0        0      0 S  0.0   0.0   0:00.00  kauditd
   23 root      20   0   0        0      0 S  0.0   0.0   0:00.00  khungtaskd
   24 root      20   0   0        0      0 S  0.0   0.0   0:00.00  oom_reaper
   27 root      0 -20   0        0      0 S  0.1   0.0   0:00.00  writeback
   28 root      20   0   0        0      0 S  0.0   0.0   0:00.02  kcompactd0
   29 root      39  19   0        0      0 S  0.0   0.0   0:00.00  kdumpd
   30 root      0 -20   0        0      0 S  0.1   0.0   0:00.00  kintegrityd
   31 root      0 -20   0        0      0 S  0.1   0.0   0:00.00  khlockd
   32 root      0 -20   0        0      0 S  0.1   0.0   0:00.00  hikup_mount_bis
   33 root      20   0   0        0      0 S  0.0   0.0   0:00.00  xen-balloon
   34 root      0 -20   0        0      0 S  0.1   0.0   0:00.00  tpm_dev_wq
   35 root      0 -20   0        0      0 S  0.1   0.0   0:00.00  md
   36 root      0 -20   0        0      0 S  0.0   0.0   0:00.00  edac-poller
   37 root      -51  0   0        0      0 S  0.0   0.0   0:00.00  watchdogd
   38 root      0 -20   0        0      0 S  0.0   0.0   0:00.01  kworker/0:1H-khlockd
   73 root      20   0   0        0      0 S  0.0   0.0   0:00.00  kswapd0
   75 root      0 -20   0        0      0 S  0.1   0.0   0:00.00  xfsailon
```

Step 16: Install stress command by using the command: “yum installstress” and type “stress -c 15” this will increase the load on cpu and makes to launch few more web servers to handle.



Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public
-	i-0ea9b93a34c10dd82	Running	t2.micro	2/2 checks passed	No alarms	ap-south-1a	ec2-5:
-	i-05340bfc565a6de77	Running	t2.micro	2/2 checks passed	No alarms	ap-south-1a	ec2-4:
-	i-096d27b95d1c194c6	Running	t2.micro	2/2 checks passed	No alarms	ap-south-1b	ec2-3:
-	i-0a1462c6df7543c87	Running	t2.micro	2/2 checks passed	No alarms	ap-south-1b	ec2-3:

-You can see in my instances that few more instances are added at us-east 1a or us-east 1b locations right after I increased the stress on cpu.

- Now you can decrease the stress by using the same command: "stress-c 1" and observe that the instances will terminate themselves.

This is the Demonstration of Auto-Scaling.